

Assignment No: 03

Title: Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months.

Name: Vasant Kumar

Class :B.E

Div:B

Batch:C

Roll No:405B091

```
In [1]: # Importing header files and csv file
#Importing the libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
# Loading file
df = pd.read_csv("Churn_Modelling.csv")
```

Preprocessing

```
In [2]: df.head()
```

```
Out[2]:   RowNumber CustomerId Surname CreditScore Geography Gender Age Tenure Balance NumOfProd
0           1  15634602 Hargrave        619    France Female  42      2     0.00
1           2  15647311    Hill         608    Spain Female  41      1  83807.86
2           3  15619304    Onio          502    France Female  42      8  159660.80
3           4  15701354    Boni          699    France Female  39      1     0.00
4           5  15737888  Mitchell        850    Spain Female  43      2  125510.82
```

◀ ▶

```
In [3]: df.shape
```

```
Out[3]: (10000, 14)
```

```
In [4]: df.describe()
```

```
Out[4]:   RowNumber CustomerId CreditScore Age Tenure Balance NumOfProducts
count  10000.00000  1.000000e+04 10000.000000 10000.000000 10000.000000 10000.000000 10000.000000
mean   5000.50000  1.569094e+07  650.528800 38.921800  5.012800  76485.889288  1.530200
std    2886.89568  7.193619e+04  96.653299 10.487806  2.892174  62397.405202  0.581654
min    1.00000  1.556570e+07  350.000000 18.000000  0.000000  0.000000  1.000000
25%   2500.75000  1.562853e+07  584.000000 32.000000  3.000000  0.000000  1.000000
50%   5000.50000  1.569074e+07  652.000000 37.000000  5.000000  97198.540000  1.000000
75%   7500.25000  1.575323e+07  718.000000 44.000000  7.000000 127644.240000  2.000000
max   10000.00000 1.581569e+07  850.000000 92.000000 10.000000 250898.090000  4.000000
```

◀ ▶

```
In [5]: df.isnull()
```

```
Out[5]:   RowNumber CustomerId Surname CreditScore Geography Gender Age Tenure Balance NumOfPr
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfPr
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False
...
9995	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False

10000 rows × 14 columns



In [6]: `df.isnull().sum()`

Out[6]:

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0
dtype: int64	

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   RowNumber        10000 non-null   int64  
 1   CustomerId      10000 non-null   int64  
 2   Surname          10000 non-null   object  
 3   CreditScore      10000 non-null   int64  
 4   Geography         10000 non-null   object  
 5   Gender            10000 non-null   object  
 6   Age               10000 non-null   int64  
 7   Tenure            10000 non-null   int64  
 8   Balance           10000 non-null   float64 
 9   NumOfProducts     10000 non-null   int64  
 10  HasCrCard         10000 non-null   int64  
 11  IsActiveMember    10000 non-null   int64  
 12  EstimatedSalary   10000 non-null   float64 
 13  Exited            10000 non-null   int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In [8]: `df.dtypes`

```
Out[8]: RowNumber      int64
CustomerId       int64
Surname         object
CreditScore      int64
Geography        object
Gender           object
Age              int64
Tenure           int64
Balance          float64
NumOfProducts    int64
HasCrCard        int64
IsActiveMember   int64
EstimatedSalary  float64
Exited           int64
dtype: object
```

```
In [9]: df.columns
```

```
Out[9]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
   'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
   'IsActiveMember', 'EstimatedSalary', 'Exited'],
   dtype='object')
```

```
In [10]: df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis= 1) #Dropping the unnecessary columns
df.head()
```

```
Out[10]: CreditScore Geography Gender Age Tenure Balance NumOfProducts HasCrCard IsActiveMember Es
0 619 France Female 42 2 0.00 1 1 1
1 608 Spain Female 41 1 83807.86 1 0 1
2 502 France Female 42 8 159660.80 3 1 0
3 699 France Female 39 1 0.00 2 0 0
4 850 Spain Female 43 2 125510.82 1 1 1
```

Visualization

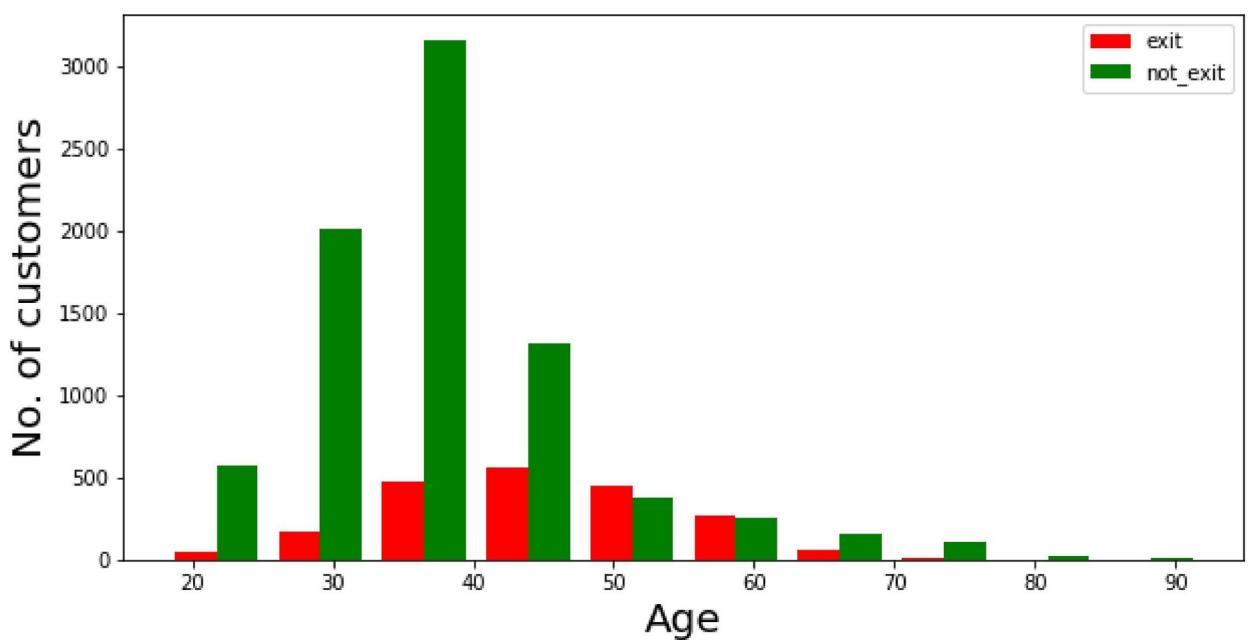
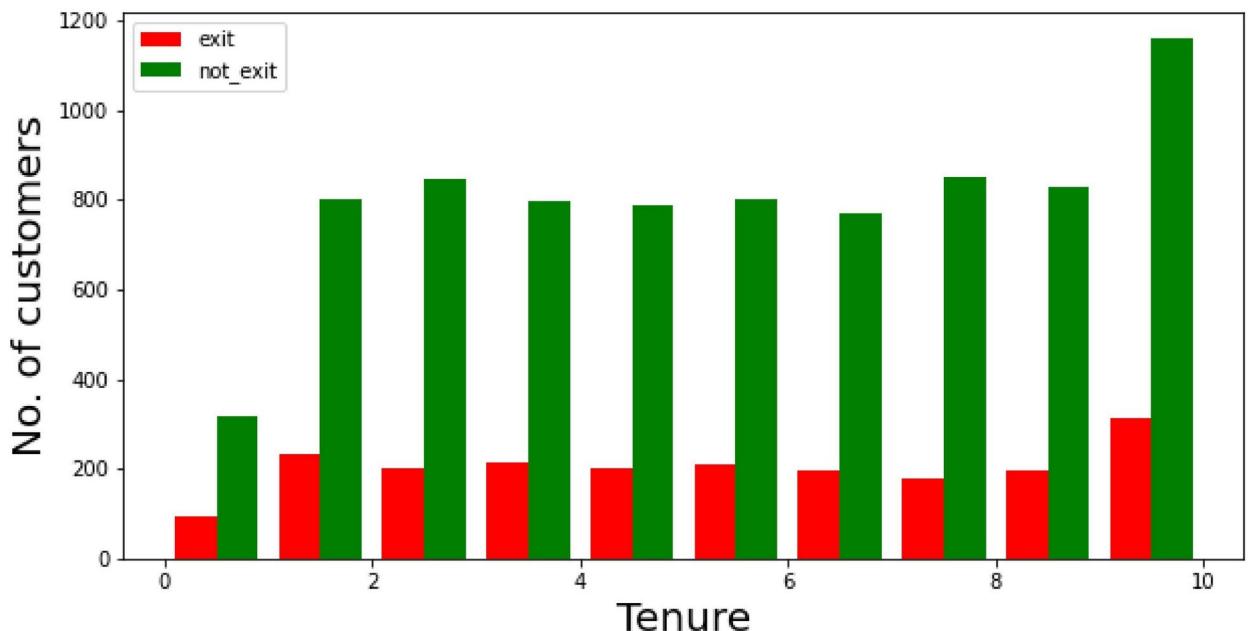
```
In [11]: def visualization(x, y, xlabel):
    plt.figure(figsize=(10,5))
    plt.hist([x, y], color=['red', 'green'], label = ['exit', 'not_exit'])
    plt.xlabel(xlabel, fontsize=20)
    plt.ylabel("No. of customers", fontsize=20)
    plt.legend()
    df_churn_exited = df[df['Exited']==1]['Tenure']
    df_churn_not_exited = df[df['Exited']==0]['Tenure']
    visualization(df_churn_exited, df_churn_not_exited, "Tenure")
    df_churn_exited2 = df[df['Exited']==1]['Age']
    df_churn_not_exited2 = df[df['Exited']==0]['Age']
    visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```

```
/usr/local/lib/python3.7/dist-packages/numpy/core/fromnumeric.py:3208: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.
```

```
    return asarray(a).size
```

```
/usr/local/lib/python3.7/dist-packages/matplotlib/cbook/__init__.py:1376: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.
```

```
    X = np.atleast_1d(X.T if isinstance(X, np.ndarray) else np.asarray(X))
```



```
In [12]: # Converting the Categorical Variables
X = df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember']]
states = pd.get_dummies(df['Geography'], drop_first = True)
gender = pd.get_dummies(df['Gender'], drop_first = True)
df = pd.concat([df, gender, states], axis = 1)
df.head()
```

```
Out[12]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	Es
0	619	France	Female	42	2	0.00		1	1	1
1	608	Spain	Female	41	1	83807.86		1	0	1
2	502	France	Female	42	8	159660.80		3	1	0
3	699	France	Female	39	1	0.00		2	0	0
4	850	Spain	Female	43	2	125510.82		1	1	1

```
In [13]: # Splitting the training and testing Dataset
X = df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'Exited']]
y = df['Exited']
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.30)
```

```
In [14]: # Normalizing the values with mean as 0 and Standard Deviation as 1
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
print("Train \n",X_train)
print("Test \n",X_test)
```

```
Train
[[ 0.62819435 -0.9478278  1.04124445 ... -1.08847096 -0.57031183
-0.57339125]
[ 0.349612   1.83232831 -1.36892977 ...  0.91871996 -0.57031183
-0.57339125]
[ 0.32897627 -0.85196035  1.72986565 ...  0.91871996 -0.57031183
-0.57339125]
...
[ 1.15440545  1.35299105 -0.33599796 ... -1.08847096 -0.57031183
-0.57339125]
[-1.86872892 -0.37262309 -1.36892977 ...  0.91871996 -0.57031183
-0.57339125]
[ 0.48374424  1.2571236   1.04124445 ...  0.91871996  1.75342671
-0.57339125]]
Test
[[ 0.80359805  0.58605143  0.69693385 ... -1.08847096 -0.57031183
1.74400987]
[ 0.41151919 -1.71476742 -0.68030856 ...  0.91871996 -0.57031183
1.74400987]
[ 0.61787649  0.20258162 -0.68030856 ...  0.91871996 -0.57031183
-0.57339125]
...
[-0.10437405 -0.18088819  0.69693385 ... -1.08847096  1.75342671
-0.57339125]
[-1.7552324   0.10671417 -0.68030856 ...  0.91871996 -0.57031183
-0.57339125]
[-0.03214899 -0.08502073  1.72986565 ...  0.91871996 -0.57031183
-0.57339125]]
```

```
In [15]: # Building the Classifier Model using Keras
import keras
from keras.models import Sequential #To create sequential neural network
from keras.layers import Dense #To create hidden layers
```

```
In [16]: classifier = Sequential()
classifier.add(Dense(activation = "relu",input_dim = 11,units = 6,kernel_initializer = "uniform"))
classifier.add(Dense(activation = "relu",units = 6,kernel_initializer = "uniform")) #Adding second layer
classifier.add(Dense(activation = "sigmoid",units = 1,kernel_initializer = "uniform")) #Final layer
classifier.compile(optimizer="adam",loss = 'binary_crossentropy',metrics = [ 'accuracy']) #To calculate accuracy
```

```
In [17]: classifier.summary() #3 Layers created. 6 neurons in 1st,6neurons in 2nd Layer and 1 neuron in output layer
```

```
Model: "sequential"

Layer (type)          Output Shape       Param #
=====
dense (Dense)         (None, 6)           72
dense_1 (Dense)       (None, 6)           42
dense_2 (Dense)       (None, 1)            7
=====
Total params: 121
Trainable params: 121
Non-trainable params: 0
```

```
In [18]: classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the ANN to training dataset  
y_pred = classifier.predict(X_test)  
y_pred = (y_pred > 0.5) #Predicting the result
```

Epoch 1/50
700/700 [=====] - 2s 1ms/step - loss: 0.5023 - accuracy: 0.7934
Epoch 2/50
700/700 [=====] - 1s 1ms/step - loss: 0.4347 - accuracy: 0.7931
Epoch 3/50
700/700 [=====] - 1s 1ms/step - loss: 0.4309 - accuracy: 0.7931
Epoch 4/50
700/700 [=====] - 1s 1ms/step - loss: 0.4267 - accuracy: 0.8071
Epoch 5/50
700/700 [=====] - 1s 1ms/step - loss: 0.4230 - accuracy: 0.8174
Epoch 6/50
700/700 [=====] - 1s 1ms/step - loss: 0.4187 - accuracy: 0.8246
Epoch 7/50
700/700 [=====] - 1s 1ms/step - loss: 0.4152 - accuracy: 0.8287
Epoch 8/50
700/700 [=====] - 1s 1ms/step - loss: 0.4124 - accuracy: 0.8304
Epoch 9/50
700/700 [=====] - 1s 1ms/step - loss: 0.4102 - accuracy: 0.8291
Epoch 10/50
700/700 [=====] - 1s 1ms/step - loss: 0.4086 - accuracy: 0.8313
Epoch 11/50
700/700 [=====] - 1s 1ms/step - loss: 0.4070 - accuracy: 0.8320
Epoch 12/50
700/700 [=====] - 1s 1ms/step - loss: 0.4059 - accuracy: 0.8317
Epoch 13/50
700/700 [=====] - 1s 1ms/step - loss: 0.4050 - accuracy: 0.8317
Epoch 14/50
700/700 [=====] - 1s 1ms/step - loss: 0.4039 - accuracy: 0.8330
Epoch 15/50
700/700 [=====] - 1s 1ms/step - loss: 0.4033 - accuracy: 0.8309
Epoch 16/50
700/700 [=====] - 1s 2ms/step - loss: 0.4018 - accuracy: 0.8326
Epoch 17/50
700/700 [=====] - 1s 2ms/step - loss: 0.4026 - accuracy: 0.8317
Epoch 18/50
700/700 [=====] - 1s 2ms/step - loss: 0.4019 - accuracy: 0.8321
Epoch 19/50
700/700 [=====] - 1s 1ms/step - loss: 0.4016 - accuracy: 0.8337
Epoch 20/50
700/700 [=====] - 1s 1ms/step - loss: 0.4012 - accuracy: 0.8320
Epoch 21/50
700/700 [=====] - 1s 1ms/step - loss: 0.4010 - accuracy: 0.8341
Epoch 22/50
700/700 [=====] - 1s 1ms/step - loss: 0.4006 - accuracy: 0.8334
Epoch 23/50
700/700 [=====] - 1s 1ms/step - loss: 0.4012 - accuracy: 0.8321
Epoch 24/50
700/700 [=====] - 1s 1ms/step - loss: 0.4003 - accuracy: 0.8341
Epoch 25/50
700/700 [=====] - 1s 1ms/step - loss: 0.4001 - accuracy: 0.8324
Epoch 26/50
700/700 [=====] - 1s 998us/step - loss: 0.3994 - accuracy: 0.8336
Epoch 27/50
700/700 [=====] - 1s 1ms/step - loss: 0.3999 - accuracy: 0.8343
Epoch 28/50
700/700 [=====] - 1s 1ms/step - loss: 0.3997 - accuracy: 0.8324
Epoch 29/50
700/700 [=====] - 1s 1ms/step - loss: 0.3990 - accuracy: 0.8334
Epoch 30/50
700/700 [=====] - 1s 993us/step - loss: 0.3985 - accuracy: 0.8351
Epoch 31/50
700/700 [=====] - 1s 1ms/step - loss: 0.3989 - accuracy: 0.8339
Epoch 32/50
700/700 [=====] - 1s 1ms/step - loss: 0.3991 - accuracy: 0.8347
Epoch 33/50
700/700 [=====] - 1s 1ms/step - loss: 0.3987 - accuracy: 0.8346
Epoch 34/50
700/700 [=====] - 1s 1ms/step - loss: 0.3983 - accuracy: 0.8349

```
Epoch 35/50
700/700 [=====] - 1s 1ms/step - loss: 0.3985 - accuracy: 0.8351
Epoch 36/50
700/700 [=====] - 1s 1ms/step - loss: 0.3984 - accuracy: 0.8349
Epoch 37/50
700/700 [=====] - 1s 2ms/step - loss: 0.3984 - accuracy: 0.8350
Epoch 38/50
700/700 [=====] - 1s 1ms/step - loss: 0.3980 - accuracy: 0.8351
Epoch 39/50
700/700 [=====] - 1s 1ms/step - loss: 0.3980 - accuracy: 0.8361
Epoch 40/50
700/700 [=====] - 1s 1ms/step - loss: 0.3984 - accuracy: 0.8331
Epoch 41/50
700/700 [=====] - 1s 1ms/step - loss: 0.3976 - accuracy: 0.8347
Epoch 42/50
700/700 [=====] - 1s 1ms/step - loss: 0.3975 - accuracy: 0.8360
Epoch 43/50
700/700 [=====] - 1s 1ms/step - loss: 0.3977 - accuracy: 0.8349
Epoch 44/50
700/700 [=====] - 1s 991us/step - loss: 0.3974 - accuracy: 0.8369
Epoch 45/50
700/700 [=====] - 1s 1ms/step - loss: 0.3972 - accuracy: 0.8347
Epoch 46/50
700/700 [=====] - 1s 1ms/step - loss: 0.3975 - accuracy: 0.8347
Epoch 47/50
700/700 [=====] - 1s 1ms/step - loss: 0.3969 - accuracy: 0.8347
Epoch 48/50
700/700 [=====] - 1s 1ms/step - loss: 0.3971 - accuracy: 0.8346
Epoch 49/50
700/700 [=====] - 1s 1ms/step - loss: 0.3969 - accuracy: 0.8353
Epoch 50/50
700/700 [=====] - 1s 1ms/step - loss: 0.3970 - accuracy: 0.8349
94/94 [=====] - 0s 853us/step
```

```
In [19]: y_pred
```

```
Out[19]: array([[False],
 [False],
 [False],
 ...,
 [False],
 [False],
 [False]])
```

```
In [20]: from sklearn.metrics import confusion_matrix,accuracy_score,classification_report
cm = confusion_matrix(y_test,y_pred)
cm
```

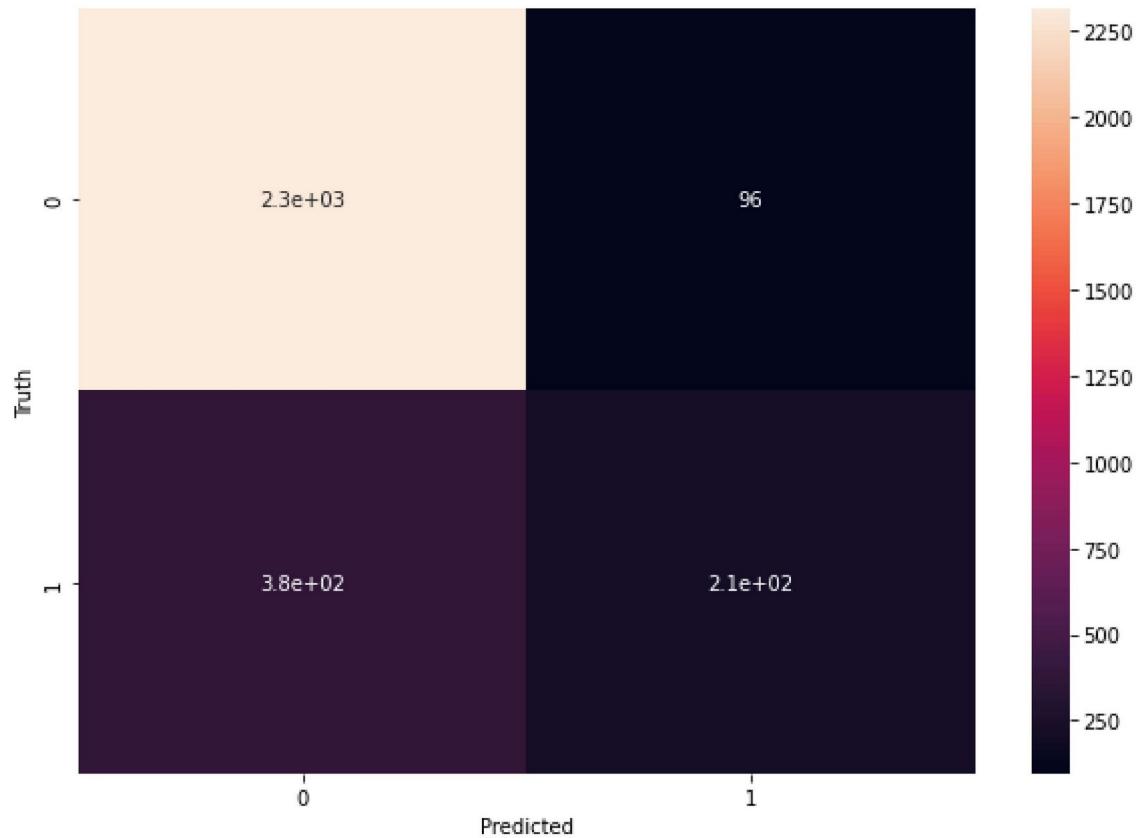
```
Out[20]: array([[2315,    96],
 [ 375,  214]])
```

```
In [21]: accuracy = accuracy_score(y_test,y_pred)
accuracy
```

```
Out[21]: 0.843
```

```
In [22]: plt.figure(figsize = (10,7))
sns.heatmap(cm,annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

```
Out[22]: Text(69.0, 0.5, 'Truth')
```



```
In [23]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.86	0.96	0.91	2411
1	0.69	0.36	0.48	589
accuracy			0.84	3000
macro avg	0.78	0.66	0.69	3000
weighted avg	0.83	0.84	0.82	3000